


Federated learning with Docker

Team 01



- 1 Introduction - What is Federated Learning?
- 2 Project Setup
 - Flower server
 - Flower client
 - Run locally
- 3 Static system
 - docker-compose.yml
 - Run the project
- 4 Dynamic System with Federated Learning
- 5 Appendix
 - Dockerfiles
 - Dynamic implementation - Bash script

What is Federated Learning?

Introduction



- Federated learning:
 - approach to machine learning
 - allows models to be trained in a *distributed way* across multiple devices or local servers
 - *no need to gather* and transfer the raw data to a central server

What is Federated Learning?

Introduction



- Federated learning:
 - approach to machine learning
 - allows models to be trained in a *distributed way* across multiple devices or local servers
 - *no need to gather* and transfer the raw data to a central server
- How does it work?
 - 1 Models are *sent to local devices* (clients)
 - 2 They are trained on local data
 - 3 Model updates are *aggregated* to form a global model

What is Federated Learning?

Introduction



- Federated learning:
 - approach to machine learning
 - allows models to be trained in a *distributed way* across multiple devices or local servers
 - *no need to gather* and transfer the raw data to a central server
- How does it work?
 - 1 Models are *sent to local devices* (clients)
 - 2 They are trained on local data
 - 3 Model updates are *aggregated* to form a global model
- Benefits:
 - protection of data *confidentiality*
 - reduced bandwidth requirements
 - ability to process geographically distributed data

What is Federated Learning?

Centralized learning




Figure: Centralized learning

What is Federated Learning?

Federated learning




Figure: Federated learning

Project Setup

Working environment



1 Clone the repository

```
1 git clone git@github.com:CCBDA-UPC/Research-Projects-2024.git
```

2 change directory to tutorial-1

```
1 cd tutorial-1/
```

3 Set up the working environment

```
1 python -m venv .venv
2 source .venv/bin/activate
3
4 pip install -r requirements.txt
```


Project Setup

Flower server



```
1 from flwr.common import ( EvaluateIns, EvaluateRes, FitIns
    , FitRes, Parameters, Scalar )
2 from flwr.server.client_manager import ClientManager
3 from flwr.server.client_proxy import ClientProxy
4 from flwr.server.strategy import Strategy
5
6
7 class FedAnalytics(Strategy):
8     def initialize_parameters(self, client_manager:
        Optional[ClientManager] = None) -> Optional[Parameters
        ]:
9         ...
10
11     def configure_fit(self, server_round: int, parameters:
        Parameters, client_manager: ClientManager) -> List[
        Tuple[ClientProxy, FitIns]]:
12         ...
13
14     def aggregate_fit( self, server_round: int, results:
        List[Tuple[ClientProxy, FitRes]], failures: List[Union
        [Tuple[ClientProxy, FitRes], BaseException]]) -> Tuple
        [Optional[Parameters], Dict[str, Scalar]]:
15         ...
```

Project Setup

Flower server



```
1  def evaluate(self, server_round: int, parameters:
    Parameters) -> Optional[Tuple[float, Dict[str, Scalar
    ]]]:
2      ...
3
4  def configure_evaluate(self, server_round: int,
    parameters: Parameters, client_manager: ClientManager)
    -> List[Tuple[ClientProxy, EvaluateIns]]:
5      ...
6
7  def aggregate_evaluate(self, server_round: int, results:
    List[Tuple[ClientProxy, EvaluateRes]], failures: List[
    Union[Tuple[ClientProxy, EvaluateRes], BaseException
    ]],) -> Tuple[Optional[float], Dict[str, Scalar]]:
8      ...
```

Project Setup

Flower server



```
1
2 # Start Flower server
3 fl.server.start_server(
4     server_address="0.0.0.0:8080",
5     config=fl.server.ServerConfig(num_rounds=num_rounds),
6     strategy=FedAnalytics(),
7 )
```

Project Setup

Flower client



```
1 from flwr_datasets import FederatedDataset
2 # Define Flower client
3 class FlowerClient(fl.client.NumPyClient):
4     def get_parameters(self, config):
5         # return the model weight as a list
6         ...
7     def set_parameters(self, parameters):
8         # update the local model weights with the
9         # parameters received from the server
10        ...
11    def fit(self, parameters, config):
12        # set the local model weights, train the model.
13        ...
14    def evaluate(self, parameters, config):
15        # test the local model
16        ...
```

Project Setup

Flower client



```
1  args = parser.parse_args()
2  partition_id = args.partition_id
3  server_address = args.server_address
4
5  # Load the partition data
6  fds = FederatedDataset(dataset="hitorilabs/iris",
7                          partitioners={"train": N_CLIENTS})
8
9  dataset = fds.load_partition(partition_id, "train").
10 with_format("pandas")[:]
11
12 # Use just the specified columns
13 X = dataset[column_names]
14
15 # Start Flower client
16 fl.client.start_client(
17     server_address=server_address,
18     client=FlowerClient(X).to_client(),
```

Project Setup

Train the model locally



Train the model (locally), federated

1 Run the server

```
1 python server.py
```

2 Run the first client

```
1 python client.py --partition-id 0 --server-  
address 127.0.0.1:8080
```

3 Run the second client

```
1 python client.py --partition-id 1 --server-  
address 127.0.0.1:8080
```

Static system

docker-compose.yml




Figure: Federated learning implementation with Docker containers



- *Portability*: consistent performance across various platforms
- *Isolation and Security*: By isolating dependencies: minimizing conflicts and potential risks
- *Ease of Deployment* and Management: simplifies the deployment and management of federated learning clusters
- *Reproducibility*: experiments can be easily shared and replicated

- The server service is built from the `server/` directory (*uses the script mentioned in the last part*)
- The server service is configured to wait for *3 rounds* before generating the global model.

```
1  server:
2      build: server/.
3      container_name: server
4      environment:
5          - NUM_ROUNDS=3
6      networks:
7          - federated_learning
```

► Dockerfile

Static system

Flower client



- The client service is built from the `client/` directory.
- The client service is configured *with 2 clients*, each with a unique partition ID and server address.

```
1  client-0:
2      build: client/.
3      restart: on-failure
4      environment:
5          - SERVER_ADDRESS=server:8080
6          - PARTITION_ID=0
7          - NUMBER_OF_CLIENTS=2
8      networks:
9          - federated_learning
```

► Dockerfile

Static system

Run the project

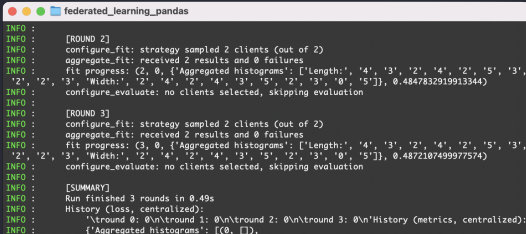


- Deploy the *docker-compose* stack {server, client-0, client-1}

```
1 docker compose up --build
```

- Check the *logs* to get the results of the process
It is possible to mount a volume and extract the stdout in a file

```
1 docker compose logs -f
```



```
federated_learning_pandas
INFO :
INFO : [ROUND 2]
INFO : configure_fit: strategy sampled 2 clients (out of 2)
INFO : aggregate_fit: received 2 results and 0 failures
INFO : fit progress: (2, 0, {'Aggregated histograms': ['Length:', '4', '3', '2', '4', '2', '5', '3',
INFO : '2', '2', '3', 'Width:', '2', '4', '2', '4', '3', '5', '2', '3', '0', '5']}, 0.4847832919913344)
INFO : configure_evaluate: no clients selected, skipping evaluation
INFO :
INFO : [ROUND 3]
INFO : configure_fit: strategy sampled 2 clients (out of 2)
INFO : aggregate_fit: received 2 results and 0 failures
INFO : fit progress: (3, 0, {'Aggregated histograms': ['Length:', '4', '3', '2', '4', '2', '5', '3',
INFO : '2', '2', '3', 'Width:', '2', '4', '2', '4', '3', '5', '2', '3', '0', '5']}, 0.4872107499977574)
INFO : configure_evaluate: no clients selected, skipping evaluation
INFO :
INFO : [SUMMARY]
INFO : Run finished 3 rounds in 0.49s
INFO : History (loss, centralized):
INFO :   \tround 0: \n\tround 1: 0 \n\tround 2: 0 \n\tround 3: 0 \n\tHistory (metrics, centralized):
INFO :   {'Aggregated histograms': [[0, []],
```

Content of *run_server.sh*

1 Create a Docker Network

```
1 docker network create federated_learning
```

2 Build the Server Docker Image

```
1 docker build -t federated_learning_server:
  latest server/.
```

3 Run the Server Docker Container

```
1 docker run --name server \
2     --env NUM_ROUNDS=8 \
3     --network federated_learning \
4     federated_learning_server
5 echo "Server started."
```

In the same way, create a Bash Script *run_random_clients.sh*

- 1 Build the image
- 2 Declare a `start_client()` function to *create one client on demand*
- 3 Declare `stop_client()` to stop a client
- 4 Main function `create_random_clients()` to *start and stop randomly clients*

```
1 ./run_random_clients.sh
```

► Bash script

Appendix

Server Dockerfile



```
1 FROM python:3.12-slim
2 WORKDIR /app
3 # Copy the requirements file
4 COPY ../requirements.txt requirements.txt
5 # Install the dependencies
6 RUN pip install -r requirements.txt --no-cache-dir
7 # Copy the rest of the code
8 COPY server.py /app
9 # Run the server
10 CMD ["python", "server.py"]
```

Appendix

Client Dockerfile



```
1 FROM python:3.12-slim
2 WORKDIR /app
3 # Copy the requirements file
4 COPY requirements.txt requirements.txt
5 # Install the dependencies
6 RUN pip install -r requirements.txt --no-cache-dir
7 # Copy the rest of the code
8 COPY client.py /app
9
10 # Add the argument
11 ENV PARTITION_ID=0
12 ENV SERVER_ADDRESS="server:8080"
13
14 # Run the server
15 CMD python client.py --partition-id $PARTITION_ID --
    server-address $SERVER_ADDRESS
```

Appendix

Run clients - 1



```
1 #!/bin/bash
2 docker build -t federated_learning_client:latest
   client/.
3
4 start_client() {
5     client_id=$1
6     partition_id=$2
7     client_name="client-${client_id}"
8     docker run -d --name "$client_name" \
9         --env SERVER_ADDRESS=server:8080 \
10        --env PARTITION_ID="$partition_id" \
11        --env NUMBER_OF_CLIENTS=2 \
12        --network federated_learning \
13        --restart on-failure \
14        federated_learning_client
15     echo "Started $client_name"
16 }
```


Appendix

Run clients - 2



```
1 stop_client() {
2     client_id=$1
3     client_name="client-${client_id}"
4     docker stop "$client_name" > /dev/null
5     docker rm "$client_name" > /dev/null
6     echo "Stopped $client_name"
7 }
8
9 create_random_clients() {
10     for ((i=0; i<10; i++)); do
11         client_id=$((RANDOM % 1000))
12         start_client "$client_id" "$i"
13         sleep 1
14     done
15 }
16
17 create_random_clients
```